

Quantifying and Reducing Stereotypes in Word Embeddings

Tolga Bolukbasi¹
 Kai-Wei Chang²
 James Zou²
 Venkatesh Saligrama¹
 Adam Kalai²

TOLGAB@BU.EDU
 KW@KWCHANG.NET
 JAMESYZOU@GMAIL.COM
 SRV@BU.EDU
 ADAM.KALAI@MICROSOFT.COM

¹ Boston University, 8 Saint Mary's Street, Boston, MA

² Microsoft Research New England, 1 Memorial Drive, Cambridge, MA

Abstract

Machine learning algorithms are optimized to model statistical properties of the training data. If the input data reflects stereotypes and biases of the broader society, then the output of the learning algorithm also captures these stereotypes. In this paper, we initiate the study of gender stereotypes in word embedding, a popular framework to represent text data. As their use becomes increasingly common, applications can inadvertently amplify unwanted stereotypes. We show across multiple datasets that the embeddings contain significant gender stereotypes, especially with regard to professions. We created a novel gender analogy task and combined it with crowdsourcing to systematically quantify the gender bias in a given embedding. We developed an efficient algorithm that reduces gender stereotype using just a handful of training examples while preserving the useful geometric properties of the embedding. We evaluated our algorithm on several metrics. While we focus on male/female stereotypes, our framework may be applicable to other types of embedding biases.

1. Introduction

Word embeddings, trained only on word co-occurrence in text corpora, capture rich semantic information about words and their meanings (Mikolov et al., 2013b). Each word (or common phrase) $w \in W$ is encoded as a d -dimensional word vector $v_w \in \mathbb{R}^d$. Using simple vector arithmetic, the embeddings are capable of answering analogy puzzles. For instance, *man:king :: woman:___*¹ returns *queen* as the answer, and similarly Japan is returned

¹An analogy puzzle, $a:b :: c:d$, involves selecting the most appropriate d given a , b , and c .

for *Paris:France :: Tokyo:Japan* (computer-generated answers are underlined). A number of such embeddings have been made publicly available including the popular word2vec (Mikolov et al., 2013a; Mikolov et al.) embedding trained on 3 million words into 300 dimensions, which we refer to here as the w2vNEWS embedding because it was trained on a corpus of text from Google News. These word embeddings have been used in a variety of downstream applications (e.g., document ranking (Nalisnick et al., 2016), sentiment analysis (İrsoy & Cardie, 2014), and question retrieval (Lei et al., 2016)).

While word-embeddings encode semantic information they also exhibit hidden biases inherent in the dataset they are trained on. For instance, word embeddings based on w2vNEWS can return biased solutions to analogy puzzles such as *father:doctor :: mother:nurse* and *man:computer programmer :: woman:homemaker*. Other publicly available embeddings produce similar results exhibiting gender stereotypes. Moreover, the closest word to the query *BLACK MALE* returns *ASSAULTED* while the response to *WHITE MALE* is *ENTITLED TO*. This raises serious concerns about their widespread use.

The prejudices and stereotypes in these embeddings reflect biases implicit in the data on which they were trained. The embedding of a word is typically optimized to predict co-occurring words in the corpus. Therefore, if *mother* and *nurse* frequently co-occur, then the vectors v_{mother} and v_{nurse} also tend to be more similar and encode the gender stereotypes. The use of embeddings in applications can amplify these biases. To illustrate this point, consider Web search where, for example, one recent project has shown that, when carefully combined with existing approaches, word vectors can significantly improve Web page relevance results (Nalisnick et al., 2016) (note that this work is a proof of concept – we do not know which, if any, mainstream search engines presently incorporate word embeddings). Consider a researcher seeking a summer intern to work on a machine learning project on deep learning who searches for, say, “linkedin graduate student machine learning neural networks.” Now, a word embedding’s semantic knowledge

can improve relevance in the sense that a LinkedIn web page containing terms such as “PhD student,” “embeddings,” and “deep learning,” which are related to but different from the query terms, may be ranked highly in the results. However, word embeddings also rank CS research related terms closer to male names than female names. The consequence would be, between two pages that differed in the names Mary and John but were otherwise identical, the search engine would rank John’s higher than Mary. In this hypothetical example, the usage of word embedding makes it even harder for women to be recognized as computer scientists and would contribute to widening the existing gender gap in computer science. While we focus on gender bias, specifically male/female, our approach may be applied to other types of biases.

We propose two methods to systematically quantify the gender bias in a set of word embeddings. First, we quantify how words, such as those corresponding to professions, are distributed along the direction between embeddings of *he* and *she*. Second, we design an algorithm for generating analogy pairs from an embedding given two seed words and we use crowdworkers to quantify whether these embedding analogies reflect stereotypes. Some analogies reflect stereotypes such as *he:janitor :: she:housekeeper* and *he:alcoholism :: she:eating disorders*. Finally, others may provoke interesting discussions such as *he:realist :: she:feminist* and *he:injured :: she:victim*.

Since biases are cultural, we enlist U.S.-based crowdworkers to identify analogies to judge whether analogies: (a) reflect *stereotypes* (to understand biases), or (b) are nonsensical (to ensure accuracy). We first establish that biases indeed exist in the embeddings. We then show that, surprisingly, information to distinguish stereotypical associations like female:homemaker from definitional associations like female:sister can often be removed. We propose an approach that, given an embedding and only a handful of words, can reduce the amount of bias present in that embedding without significantly reducing its performance on other benchmarks.

Contributions. (1) We initiate the study of stereotypes and biases in word embeddings. Our work follows a large body of literature on bias in language, but word embeddings are of specific interest because they are commonly used in machine learning and they have simple geometric structures that can be quantified mathematically. (2) We develop two metrics to quantify gender stereotypes in word embeddings based on words associated with professions together with automatically generated analogies which are then scored by the crowd. (3) We develop a new algorithm that reduces gender stereotypes in the embedding using only a handful of training examples while preserving useful properties of the embedding.




Figure 1. Comparison of gender bias of profession words across two embeddings: word2vec trained on Googlenews and GloVe trained web-crawl texts. The x and y axes show projections onto the *he-she* direction in the two embeddings. Each dot is one of 249 common profession words. Words closest to *he*, closest to *she*, and in between the two are colored in red and shown in the plot.

Prior work. The body of prior work on bias in language and prejudice in machine learning algorithms is too large to fully cover here. We note that gender stereotypes have been shown to develop in children as young as two years old (Turner & Gervai, 1995). Statistical analyses of language have shown interesting contrasts between language used to describe men and women, e.g., in recommendation letters (Schmader et al., 2007). A number of online systems have been shown to exhibit various biases, such as racial discrimination in the ads presented to users (Sweeney, 2013). Approaches to modify classification algorithms to define and achieve various notions of fairness have been described in a number of works, see, e.g., (Barocas & Selbst, 2014; Dwork et al., 2012) and a recent survey (Zliobaite, 2015).

2. Implicit stereotypes in word embedding

Stereotyped words. A simple approach to explore how gender stereotypes manifest in embeddings is to quantify which words are closer to *he* versus *she* in the embedding space (using other words to capture gender, such as *man* and *woman*, gives similar but noisier results due to their multiple meanings). We used a list of 215 common profession names, removing names that are associated with one gender by definition (e.g. waitress, waiter). For each name, v , we computed its projection onto the gender axis: $v \cdot (v_{he} - v_{she}) / \|v_{he} - v_{she}\|_2$. Figure 1 shows the projection of professions on the w2vNEWS embedding (x -axis) and on a different embedding trained by GloVe on a dataset of web-crawled texts (y -axis). Several professions are closer to the *he* or *she* vector and this is consistent

across the embeddings, suggesting that embeddings encode gender stereotypes.

Stereotyped analogies. While professions give easily-interpretable insights on embedding stereotypes, we developed a more general method to automatically detect and quantify gender bias in any word embedding. Embeddings have shown to perform well in analogy tasks. Motivated by this, we ask the embedding to generate analogous word pairs for *he* and *she*, and use crowd-sourcing to evaluate the degree of stereotype of each pair.

A desired analogy $he:she :: w_1:w_2$ has the following properties²: 1) the direction of w_1-w_2 has to align with $he-she$; 2) w_1 and w_2 should be semantically similar, i.e. $\|w_1 - w_2\|_2$ is not too large. Based on this, given a word embedding E , we proposed to score analogous pairs by the following formulation:

$$S_d(w_a, w_b) = \frac{(w_a - w_b) \cdot d}{\|w_a - w_b\|_2} \text{ s.t. } \|w_a - w_b\|_2 \leq \delta \quad (1)$$

where $d = (v_{he} - v_{she}) / \|v_{he} - v_{she}\|_2$ is the gender direction and δ is a threshold for similarity.³ We observe that setting $\delta = 1$ often works well in practice; this corresponds to requiring that the two words forming the analogy are significantly closer together than two random embedding vectors.

From the embedding, we generated the top analogous pairs with the largest S_d scores. To avoid redundancies, if multiple pairs share the same w_a or w_b , we kept only one pair. Then we employed Amazon Mechanical Turk to evaluate the analogies. For each analogy, such as *man:woman :: doctor:nurse*, we ask the Turkers two yes/no questions to verify if this pairing makes sense as an analogy and whether it exhibits gender stereotype. Every word pair is judged by 10 Turkers, and we used the number of Turkers that rated this pair as stereotyped to quantify the degree of bias of this analogy. Table 1 shows the most and least stereotypical analogies generated by word2vec on Googlenews. Overall, 21% and 32% analogy judgments were stereotypical and nonsensical, respectively, by the Turkers.

3. Reducing stereotypes in word embedding

Having demonstrated that word embeddings contain substantial stereotypes in both professions and analogies, we developed a method to reduce these stereotypes while preserving the desirable geometry of the embedding.

Word embeddings are often trained on a large corpus

²For the ease of presentation, we abuse the notation to use w to represent a word or a word vector depending on the context.

³We explored alternatives including a variation of 3-CosMul (Levy & Goldberg, 2014) for generating word pairs, and observe that the proposed approach works the best.

(w2vNEWS is trained on Google news corpus with 100 billion words). As a result, it is impractical and even impossible (the corpus is not publicly accessible) to reduce the stereotypes during the training of the word vectors. Therefore, we assume that we are given a set of word vectors and aim to remove stereotypes as a post-processing step.

Our approach takes the following as inputs: **(1)** a word embedding stored in a matrix $E \in \mathbb{R}^{n,r}$, where n is the number of words and r is the dimension of the latent space. **(2)** A matrix $B \in \mathbb{R}^{n_b,r}$ where each column is a vector representing a direction of stereotype. In this paper, $B = v_{he} - v_{she}$, but in general, B can contain multiple stereotypes including gender, racism, etc.⁴ **(3)** A matrix $P \in \mathbb{R}^{n_p,r}$ whose columns correspond to set of seed words that we want to debias. An example of a seed word for gender is *manager*. **(4)** A matrix $A \subseteq E$ whose columns represent a background set of words. We want the algorithm to preserve their pairwise distances.⁵

The goal is to generate a transformation matrix $\hat{T} \in \mathbb{R}^{r,r}$, which has the following properties:

- The transformed embeddings are stereotypical-free. That is every column vectors in PT should be perpendicular to column vectors in BT (i.e., $PTT^T B^T \approx 0$).
- The transformed embeddings preserve the distances between any two vectors in the matrix A .

Let $X = TT^T$, we can capture these two objectives as the following semi-positive definite programming problem.

$$\min_{X \succeq 0} \|AXA^T - AA^T\|_F^2 + \lambda \|PXB^T\|_F^2 \quad (2)$$

where $\|\cdot\|_F$ is the Frobenius norm, the first term ensures that the pairwise distances are preserved, and the second term induces the biases to be small on the seed words. The user-specified parameter λ balances the two terms.

Directly solving this SDP optimization problem is challenging. In practice, the dimension of matrix A is in the scale of $400,000 \times 300$. The dimensions of the matrices AXA^T and AA^T are $400,000 \times 400,000$, causing computational and memory issues. We conduct singular value decomposition on A , such that $A = U\Sigma V^T$, where U and V are orthogonal matrices and Σ is a diagonal matrix.

$$\begin{aligned} \|AXA^T - AA^T\|_F^2 &= \|A(X - I)A^T\|_F^2 \\ &= \|U\Sigma V^T(X - I)V\Sigma U^T\|_F^2 \\ &= \|\Sigma V^T(X - I)V\Sigma\|_F^2. \end{aligned} \quad (3)$$

⁴Here we assume the stereotypical directions are given. In practice, this can be obtained by subjecting the vectors of the extreme words in the concept (e.g. *he* and *she* representing gender.)

⁵ Typically, we can set A to contain the word vectors in E except the ones in B and P .

Ranked as M-F stereotypical by 10/10 workers:

surgeon:nurse	doctors:midwives	athletes:gymnasts	paramedic:registered nurse
Hummer:minivan	Karate:Gymnastics	woodworking:quilting	alcoholism:eating disorders
athlete:gymnast	neurologist:therapist	hockey:figure skating	architect:interior designer
chauffeur:nanny	curator:librarian	pilots:flight attendant	drug trafficking:prostitution
musician:dancer	beers:cocktails	Sopranos:Real Housewives	headmaster:guidance counselor
workout:Pilates	Home Depot:JC Penney	weightlifting:gymnastics	Sports Illustrated:Vanity Fair
carpentry:sewing	accountant:paralegal	addiction:eating disorder	professor emeritus:associate professor

Ranked as M-F stereotypical by 0/10 workers: (random sample of 12 out of 505)

Jon:Heidi	Ainge:Fulmer	Allan:Lorna	George Clooney:Penelope Cruz
Erick:Karla	gentlemen:ladies	Christopher:Jennifer	veterans:servicemen
sausages:buns	patriarch:matriarch	Leroy:Lucille	Phillip:Belinda

Table 1. Sample of the top 1,000 analogies generated for $he:she :: w_a:w_b$ on w2vNEWS, ordered by the number of workers who judged them to reflect stereotypes. The analogies which were rated stereotypical by 10/10 workers are shown and a random sample of twelve analogies rated as stereotypical by 0/10 workers is shown. Overall, 21% of the 1000 analogies were rated as reflecting gender stereotypes.

The last equality follows the fact that U is an orthogonal matrix ($\|UXU^T\|_F^2 = \text{tr}(UXU^TUX^T U^T) = \text{tr}(UXX^T U^T) = \text{tr}(XX^T U^T U) = \|X\|_F^2$.)

Substituting Eq. (3) to Eq. (2) gives

$$\min_{X \geq 0} \|\Sigma V^T(X - I)V\Sigma\|_F^2 + \lambda \|PXB^T\|_F^2 \quad (4)$$

Here $\Sigma V^T(X - I)V\Sigma$ is a 300×300 matrix and can be solved efficiently. The solution T is the debiasing transformation of the word embedding.

Experimental validation. To validate our debiasing algorithm, we asked Turkers to suggest words that are likely to reflect gender stereotype (e.g. *manager*, *nurse*). We collected 438 such words, of which a random setup of 350 are used for training as the columns of the P matrix. The remaining are used for testing. Figure 2 illustrates the results of the algorithm. The blue circles are the 88 gender-stereotype words suggested by the Turkers which form our held-out test set. The green crosses are a random sample of background words that were not suggested to have stereotype. Most of the stereotype words lie close to the $y = 0$ line, consistent with them lies near the midpoint between *he* and *she*. In contrast the background points were substantially less affected by the debiasing transformation.

We use variances to quantify this result. For each test word (either gender-stereotypical or background) we project it onto the *he* - *she* direction. Then we compute the variance of the projections in the original embedding and after the debiasing transformation. For the gender-stereotype test words, the variance in the original embedding is 0.02 and the variance after the transformation is 0.001. For the background words, the variance before and after the transformation was 0.005 and 0.0055 respectively. This demonstrates that the transformation was able to reduce gender stereotype.

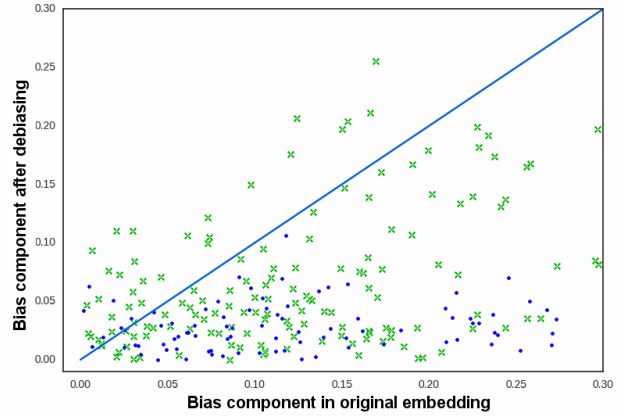


Figure 2. The changes of word vectors on the gender direction. The x and y axes show the absolute values of the projections onto the *he-she* direction before and after debiasing. The solid line is the diagonal. The blue ‘•’ are gender-stereotypical words in the test set, and the green ‘x’ are randomly selected other words that were not suggested to be gender related.

Lastly to verify that the debiasing transformation T preserves the desirable geometric structure of the embedding, we tested the transformed embedding a several standard benchmarks that measure whether related words have similar embeddings as well as how well the embedding performs in analogy tasks. Table 2 shows the results on the original and the transformed embeddings and the transformation does not negatively impact the performance.

Model	RG	WS353	RW	MSR-analogy
Before	0.761	0.700	0.471	0.712
After	0.764	0.700	0.472	0.712

Table 2. The columns show the performance of the word embeddings on the standard evaluation metrics. RG (Rubenstein & Goodenough, 1965), RW (Luong et al., 2013), WS353 (Finkelstein et al., 2001), MSR-analogy (Mikolov et al., 2013b)

References

- Barocas, Solon and Selbst, Andrew D. Big data’s disparate impact. *Available at SSRN 2477899*, 2014.
- Dwork, Cynthia, Hardt, Moritz, Pitassi, Toniann, Reingold, Omer, and Zemel, Richard. Fairness through awareness. In *Innovations in Theoretical Computer Science Conference*, 2012.
- Finkelstein, Lev, Gabrilovich, Evgeniy, Matias, Yossi, Rivlin, Ehud, Solan, Zach, Wolfman, Gadi, and Ruppín, Eytan. Placing search in context: The concept revisited. In *WWW*. ACM, 2001.
- İrsoy, Ozan and Cardie, Claire. Deep recursive neural networks for compositionality in language. In *NIPS*. 2014.
- Lei, Tao, Joshi, Hrishikesh, Barzilay, Regina, Jaakkola, Tommi, Katerina Tymoshenko, Alessandro Moschitti, and Marquez, Luis. Semi-supervised question retrieval with gated convolutions. In *NAACL*. 2016.
- Levy, Omer and Goldberg, Yoav. Linguistic regularities in sparse and explicit word representations. In *CoNLL*, 2014.
- Luong, Thang, Socher, Richard, and Manning, Christopher D. Better word representations with recursive neural networks for morphology. In *CoNLL*, pp. 104–113. Citeseer, 2013.
- Mikolov, Tomas, Sutskever, Ilya, Chen, Kai, Corrado, Gregory S., and Dean, Jeffrey. Distributed representations of words and phrases and their compositionality. In *NIPS*.
- Mikolov, Tomas, Chen, Kai, Corrado, Greg, and Dean, Jeffrey. Efficient estimation of word representations in vector space. In *ICLR*, 2013a.
- Mikolov, Tomas, Yih, Wen-tau, and Zweig, Geoffrey. Linguistic regularities in continuous space word representations. In *HLT-NAACL*, pp. 746–751, 2013b.
- Nalisnick, Eric, Mitra, Bhaskar, Craswell, Nick, and Caruana, Rich. Improving document ranking with dual word embeddings. In *www*, April 2016.
- Rubenstein, Herbert and Goodenough, John B. Contextual correlates of synonymy. *Communications of the ACM*, 8 (10):627–633, 1965.
- Schmader, Toni, Whitehead, Jessica, and Wysocki, Vicki H. A linguistic comparison of letters of recommendation for male and female chemistry and biochemistry job applicants. *Sex Roles*, 57(7-8):509–514, 2007.
- Sweeney, Latanya. Discrimination in online ad delivery. *Queue*, 11(3):10, 2013.
- Turner, Patricia J and Gervai, Judit. A multidimensional study of gender typing in preschool children and their parents: Personality, attitudes, preferences, behavior, and cultural differences. *Developmental Psychology*, 31 (5):759, 1995.
- Zliobaite, Indre. A survey on measuring indirect discrimination in machine learning. *arXiv preprint arXiv:1511.00148*, 2015.